

# Deep Belief Nets and Stacked Denoising Autoencoders

# Outline

- Neural Networks Fall
- Deep Learning's Evolution
- Deep Learning with Pre-training
- Restricted Boltzmann Machines
- Deep Belief Nets (DBNs)
- Denoising Autoencoders
- Stacked Denoising Autoencoders (SDA)
- Summary

# Neural Networks Fall

- Nonlinear problems can be learned and solved by inserting a hidden layer between the input and output layer
  - More layer, more pattern to express
- Theoretically, neural networks can approximate any function
  - Ignore time cost and over-fitting problem

# But NNs didn't Work Well

- Some cases even have less accuracy!
- Backpropagation problem
  - An error is reversed in each layer
  - The weight of the network is adjusted at each layer in order (from output to input)
  - The error gradually disappears every time it backpropagates layers
  - **Vanishing gradient problem!!**

# Outline


- Neural Networks Fall
- **Deep Learning's Evolution**
- Deep Learning with Pre-training
- Restricted Boltzmann Machines
- Deep Belief Nets (DBNs)
- Denoising Autoencoders
- Stacked Denoising Autoencoders (SDA)
- Summary

# Deep Learning's Evolution

- There are two algorithms that triggered deep learning's popularity
  - **DBN** by Hinton
    - <https://www.cs.toronto.edu/~hinton/absps/fastnc.pdf>
  - **SDA** by Vincent et al.
    - [http://www.iro.umontreal.ca/~vincentp/Publications/denoising\\_autoencoders\\_tr1316.pdf](http://www.iro.umontreal.ca/~vincentp/Publications/denoising_autoencoders_tr1316.pdf)
- So, what is the common approach that solved the vanishing gradient problem?

# Simple and Elegant Solution

- DBN and SDA use **layer-wise training** to solve vanishing gradient problem
  - Each layer adjusts the weights of the networks **independently**




# Pre-training and Fine-tuning

- This phase of layer-wise training is called **pre-training**
- The last adjustment phase is called **fine-tuning**
- The problem: if both layers are hidden (neither of the layers are input nor output layers), then how is the training done?





# Piled Up Layers in Deep Structure



# Learning Feature

- Features are learned from the input data in stages (and semi-automatically)
  - Where the deeper a layer becomes, the higher the feature it learns
  - “a machine can learn a concept”





# In NN...


- Learning intends to minimize errors between the model's prediction output and the dataset output
  - Method: remove an error by finding a pattern from the input data and making data with a common pattern the same output value (for example, 0 or 1)
- What would then happen if we **turned the output value into the input value**?
  - The weight of networks should be adjusted to focus more on the part that **reflects the common features!!**

# Pre-training

- The layer after the pre-training can be treated as normal feed-forward neural networks where the weight of the networks is adjusted
- After pre-training, features learned, and?
  - Pre-training is **unsupervised training**
  - Doesn't solve the classification problem
  - Fine-tuning to solve

# Fine-tuning


- The main roles of fine-tuning
  - Add an output layer that completed pre-training and to perform supervised training
  - Do final adjustments



# Fine-tuning

- Normally, the weights of whole networks, including the weights adjusted in pre-training, will also be adjusted
  - Deep neural networks as one multi-layer neural network
  - Doesn't the vanishing gradient problem occur?
  - Once the pre-training is done, the learning starts from the network almost already adjusted
    - Proper error can be propagated to a layer close to an input layer

# Pre-training with Fine-tuning






# Outline

- Neural Networks Fall
- Deep Learning's Evolution
- Deep Learning with Pre-training
- **Restricted Boltzmann Machines**
- Deep Belief Nets (DBNs)
- Denoising Autoencoders
- Stacked Denoising Autoencoders (SDA)
- Summary


# Restricted Boltzmann Machines

- Restricted Boltzmann Machines, RBM
- Boltzmann Machines, BM



# Boltzmann Machines, BM

- The feature is to adopt the concept of energy in neural networks
- Fully connected, take an enormous amount of calculation time



# RBM

- RBM with binary inputs is sometimes called Bernoulli RBM
- RBM is the energy-based model
  - Visible to hidden units

$$p(h_j = 1 | v) = \sigma \left( \sum_{i=1}^D w_{ij} v_i + c_j \right)$$

- Hidden to visible units

$$p(v_i = 1 | h) = \sigma \left( \sum_{j=1}^M w_{ij} h_j + b_i \right)$$

# Energy Function in RBM

$$\begin{aligned} E(v, h) &= -b^T v - c^T h - h^T W v \\ &= -\sum_{i=1}^D b_i v_i - \sum_{j=1}^M c_j h_j - \sum_{j=1}^M \sum_{i=1}^D h_j w_{ij} v_i \end{aligned}$$

- Joint probability density function

$$p(v, h) = \frac{1}{Z} \exp(-E(v, h))$$

$$Z = \sum_{v, h} \exp(-E(v, h))$$

$$p(v | \theta) = \sum_h P(v, h) = \frac{1}{Z} = \sum_h \exp(-E(v, h))$$

# Log Likelihood

$$\begin{aligned} \ln L(\theta | v) &= \ln p(v | \theta) \\ &= \ln \frac{1}{Z} \sum_h \exp(-E(v, h)) \\ &= \ln \frac{1}{Z} \sum_h \exp(-E(v, h)) - \ln \sum_{v, h} \exp(-E(v, h)) \end{aligned}$$

# Gradient

$$\begin{aligned}\frac{\partial \ln L(\theta | v)}{\partial \theta} &= \frac{\partial}{\partial \theta} \left( \ln \sum_h \exp(-E(v, h)) \right) - \frac{\partial}{\partial \theta} \left( \ln \sum_{v, h} \exp(-E(v, h)) \right) \\ &= - \frac{1}{\sum_h \exp(-E(v, h))} \sum_h \exp(-E(v, h)) \frac{\partial E(v, h)}{\partial \theta} \\ &= + \frac{1}{\sum_h \exp(-E(v, h))} \sum_{v, h} \exp(-E(v, h)) \frac{\partial E(v, h)}{\partial \theta} \\ &= - \sum_h p(h | v) \frac{\partial E(v, h)}{\partial \theta} + \sum_{v, h} p(v | h) \frac{\partial E(v, h)}{\partial \theta}\end{aligned}$$

# Gradient of each Parameter

$$\begin{aligned}\frac{\partial \ln L(\theta | v)}{\partial w_{ij}} &= \sum_h p(h | v) \frac{\partial E(v, h)}{\partial w_{ij}} + \sum_{v, h} p(h | v) \frac{\partial E(v, h)}{\partial w_{ij}} \\ &= \sum_h p(h | v) h_j v_i - \sum_v p(v) \sum_h p(h | v) h_j v_i \\ &= p(H_j = 1 | v) v_i - \sum_v p(v) p(H_j = 1 | v) v_i\end{aligned}$$

$$\frac{\partial \ln L(\theta | v)}{\partial b_i} = v_i - \sum_v p(v) v_i$$

$$\frac{\partial \ln L(\theta | v)}{\partial c_j} = p(H_j = 1 | v) - \sum_v p(v) p(H_j = 1 | v)$$



# Problem on Gradient

- Problem occurs: calculation of the probability distribution for all the  $\{0, 1\}$  patterns
  - Can't be solve within a realistic time
- Contrastive Divergence (CD)
  - The method for approximating data using Gibbs sampling

# Contrastive Divergence

Here,  $v^{(0)}$  is an input vector. Also,  $v^{(k)}$  is an input (output) vector that can be obtained by sampling for k-times using this input vector.

Then, we get:

$$h_j^{(k)} \sim p(h_j | v^{(k)})$$

$$h_i^{(k+1)} \sim p(v_i | h^{(k)})$$

$$\begin{aligned}\frac{\partial \ln L(\theta | v)}{\partial \theta} &= -\sum_h p(h | v) \frac{\partial E(v, h)}{\partial \theta} + \sum_{v, h} p(v, h) \frac{\partial E(v, h)}{\partial \theta} \\ &\approx -\sum_h p(h | v^{(0)}) \frac{\partial E(v, h)}{\partial \theta} \sum_{v, h} p(h, v^{(k)}) \frac{\partial E(v^{(k)}, h)}{\partial \theta}\end{aligned}$$

$$w_{ij}^{(\tau+1)} = w_{ij}^{(\tau)} + \eta \left( p(H_j = 1 | v^{(0)}) v_i^{(0)} - p(H_j = 1 | v^k) v_i^k \right)$$

$$b_i^{(\tau+1)} = b_i^{(\tau)} + \eta \left( v_i^{(0)} - v_i^k \right)$$

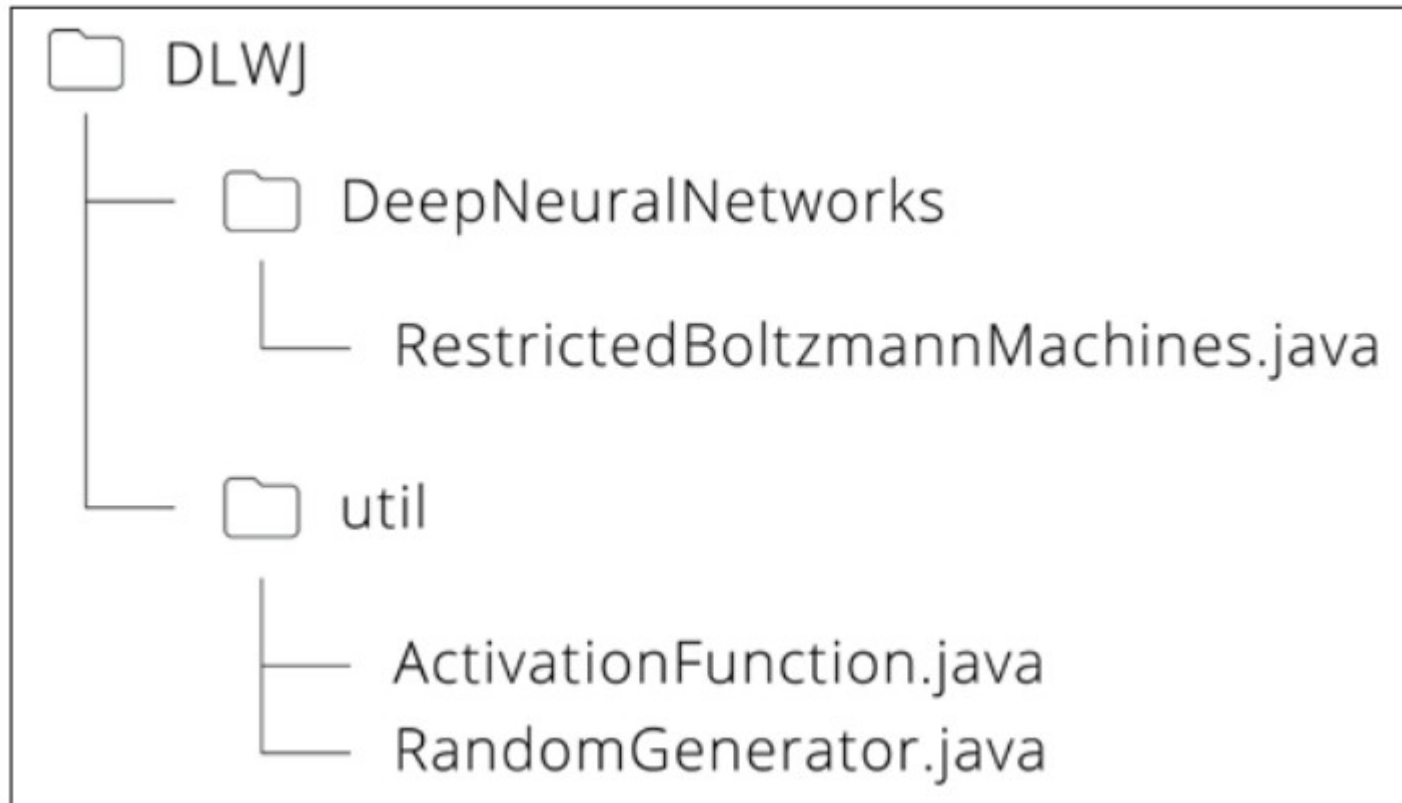
$$c_j^{(\tau+1)} = c_j^{(\tau)} + \eta \left( p(H_j = 1 | v^{(0)}) - p(H_j = 1 | v^k) \right)$$

$\tau$  is the number of iterations and  $\eta$  is the learning rate

# Contrastive Divergence

- CD that performs sampling  $k$ -times is shown as CD- $k$
- It's known that CD-1 is sufficient when applying the algorithm to realistic problems

# RBM Implementation

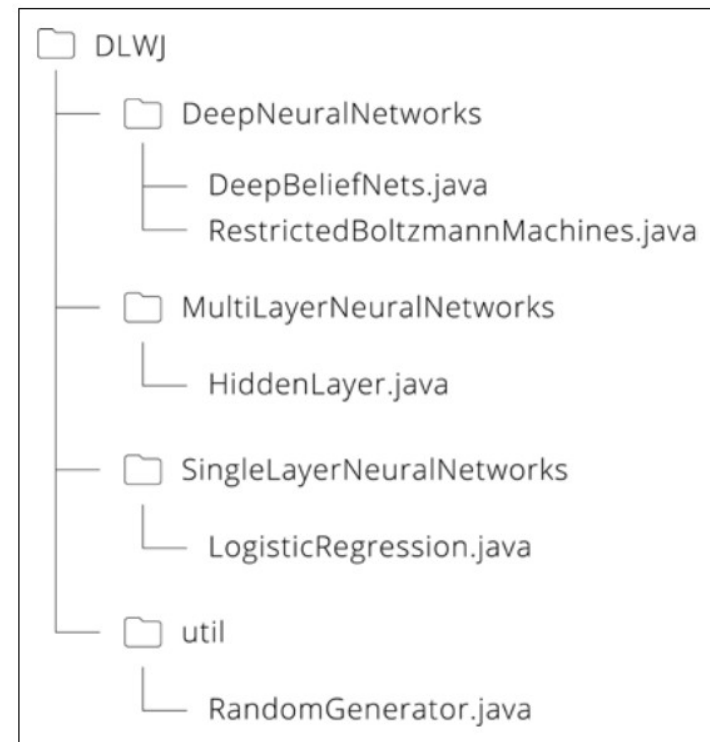


# Outline

- Neural Networks Fall
- Deep Learning's Evolution
- Deep Learning with Pre-training
- Restricted Boltzmann Machines
- **Deep Belief Nets (DBNs)**
- Denoising Autoencoders
- Stacked Denoising Autoencoders (SDA)
- Summary

# Deep Belief Nets (DBNs)

- Program flow
  - Setting up parameters for the model
  - Building the model
  - Pre-training the model
  - Fine-tuning the model
  - Testing and evaluating the model



# Outline


- Neural Networks Fall
- Deep Learning's Evolution
- Deep Learning with Pre-training
- Restricted Boltzmann Machines
- Deep Belief Nets (DBNs)
- **Denoising Autoencoders**
- Stacked Denoising Autoencoders (SDA)
- Summary



# Denoising Autoencoders

- The method used in pre-training for SDA is called Denoising Autoencoders (DA)
  - DA is the method that emphasizes the role of equating inputs and outputs
- DA processing
  - adds some noise to input data intentionally
  - Data is partially damaged
  - DA performs learning as it restores corrupted data

# Denoising Autoencoders



# Denoising Autoencoders

$$z_j = \sigma \left( \sum_{i=1}^N w_{ij} \tilde{x}_i + c_j \right)$$

$$y_i = \sigma \left( \sum_{j=1}^M w_{ji} z_j + b_i \right)$$

$\tilde{x}$  is the corrupted data, the input data with noise

# Evaluation Function of DA

$$E := -\ln L(\theta) = -\sum_{i=1}^N \left\{ x_i \ln y_i + (1 - x_i) \ln(1 - y_i) \right\}$$

- Gradient

$$h_j := \sum_{i=1}^N w_{ji} \tilde{x}_i + c_j$$

$$z_j = \sigma(h_j)$$

$$g_i := \sum_{j=1}^M w_{ji} z_j + b_i$$

$$y_i = \sigma(g_i)$$

Therefore, only two terms are required. Let's derive them one by one:

$$\frac{\partial E}{\partial h_j} = \frac{\partial E}{\partial z_j} \frac{\partial z_j}{\partial h_j} = \frac{\partial E}{\partial z_j} z_j (1 - z_j)$$

Here, we utilized the derivative of the sigmoid function:

$$\frac{d}{dx} \sigma(x) = \sigma(x)(1 - \sigma(x))$$

Also, we get:

$$\begin{aligned} \frac{\partial E}{\partial z_j} &= \sum_{i=1}^N \frac{\partial E}{\partial y_i} \frac{\partial y_i}{\partial z_j} \\ &= \sum_{i=1}^N w_{ji} (x_i - y_i) \end{aligned}$$

$$\frac{\partial E}{\partial h_j} = \left( \sum_{i=1}^N w_{ji} (x_i - y_i) \right) z_j (1 - z_j)$$

$$\begin{aligned} \frac{\partial E}{\partial g_i} &= \frac{\partial E}{\partial y_i} \frac{\partial y_i}{\partial g_i} \\ &= x_i - y_i \end{aligned}$$

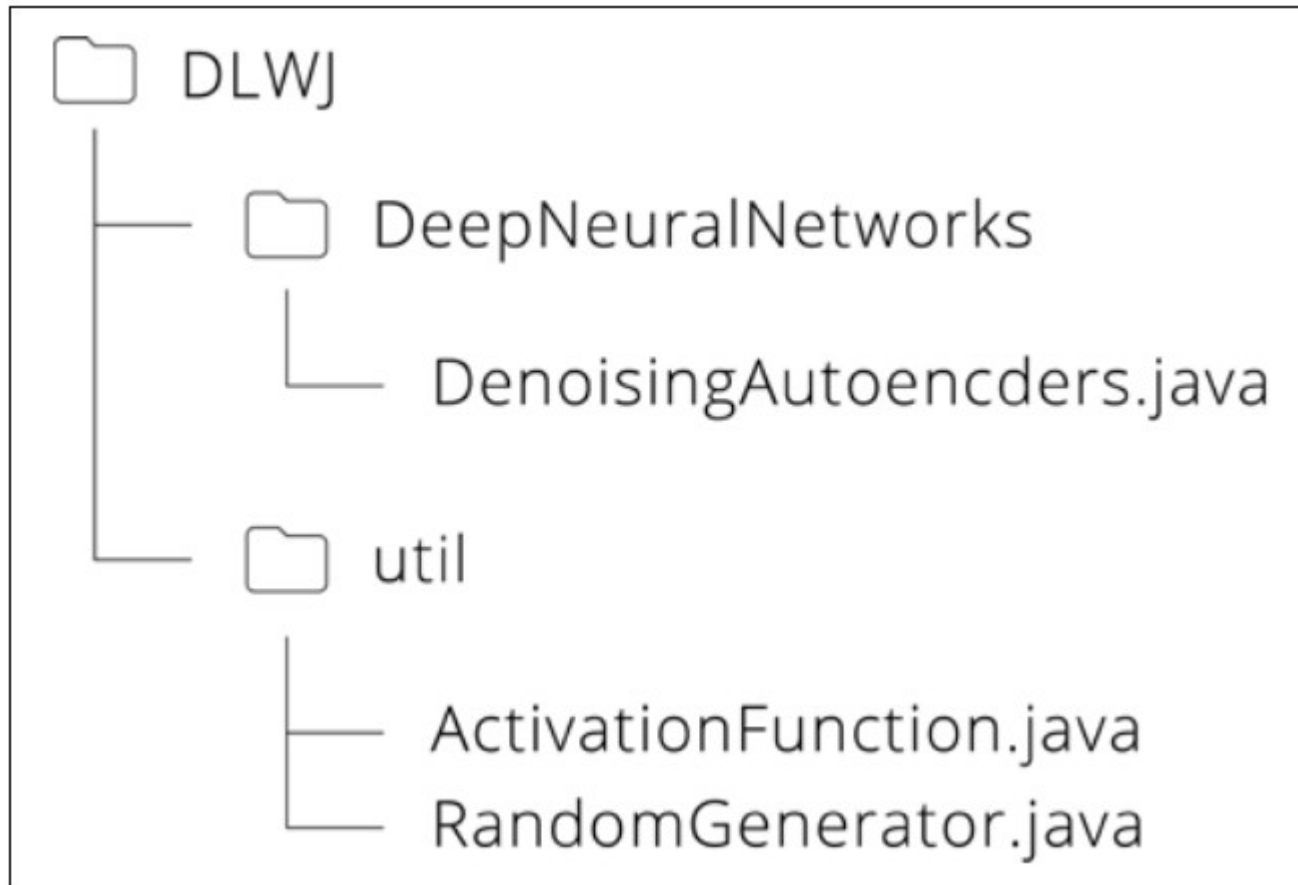
$$w_{ji}^{(k+1)} = w_{ji}^{(k)} + \eta \left[ \left( \sum_{i=1}^N w_{ji}^{(k)} (x_i - y_i) \right) z_j (1 - z_j) \tilde{x}_i + (x_i - y_i) z_j \right]$$

$$b_i^{(k+1)} = b_i^k + \eta (x_i - y_i)$$

$$c_j^{(k+1)} = c_j^{(k)} + \eta \left( \sum_{i=1}^N w_{ji}^{(k)} (x_i - y_i) \right) z_j (1 - z_j)$$

k is the number of iterations and  $\eta$  is the learning rate

# DA Implementation

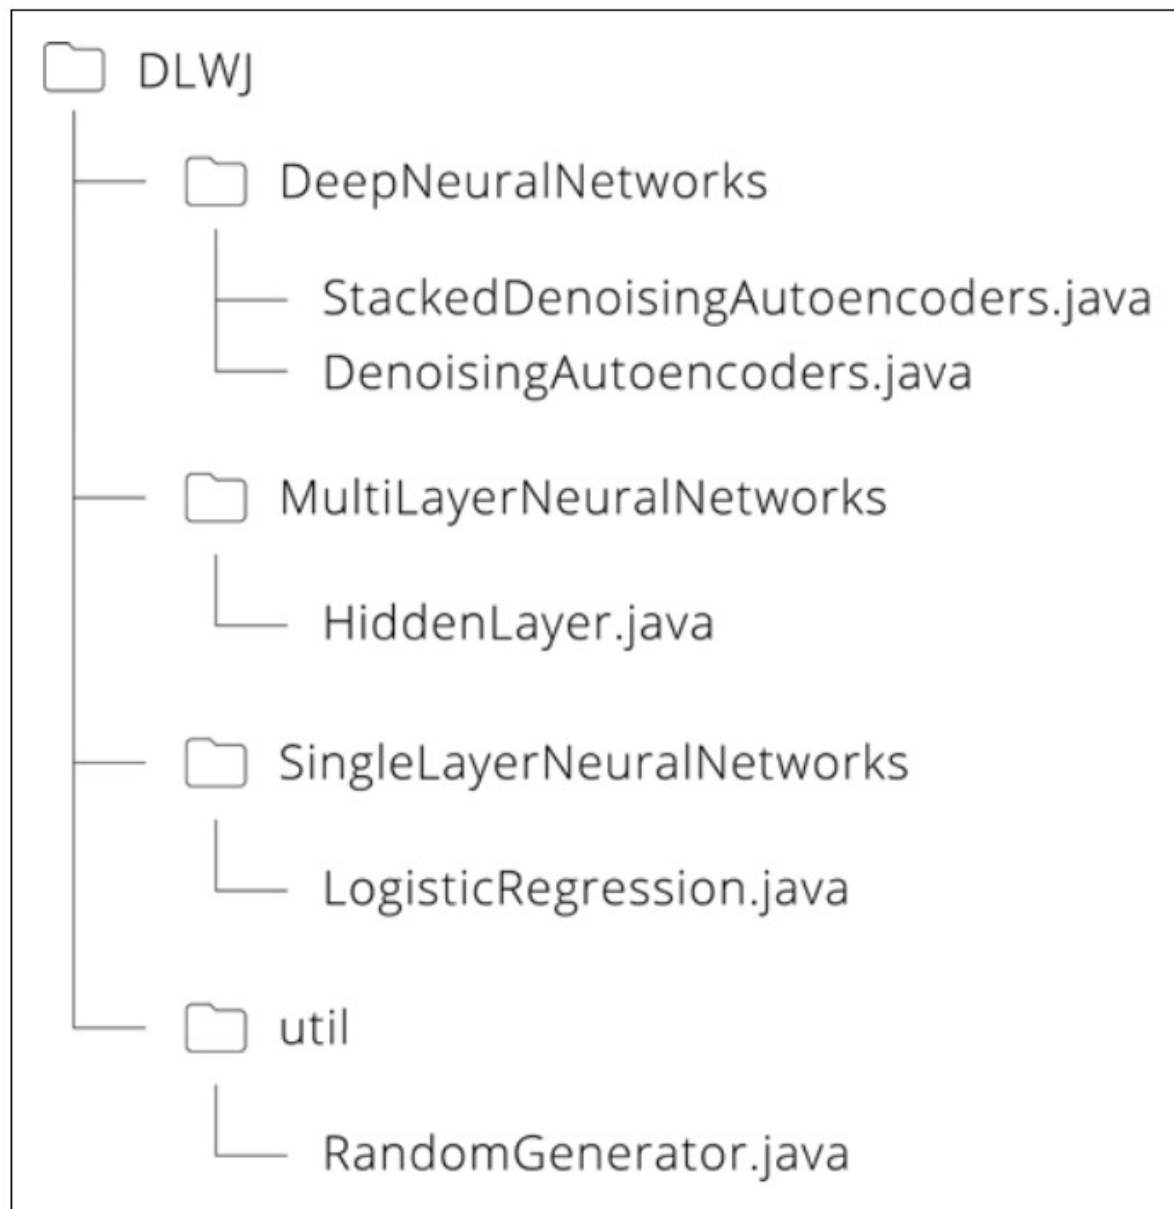


# Outline

- Neural Networks Fall
- Deep Learning's Evolution
- Deep Learning with Pre-training
- Restricted Boltzmann Machines
- Deep Belief Nets (DBNs)
- Denoising Autoencoders
- **Stacked Denoising Autoencoders (SDA)**
- Summary



# Stacked Denoising Autoencoders



# Outline

- Neural Networks Fall
- Deep Learning's Evolution
- Deep Learning with Pre-training
- Restricted Boltzmann Machines
- Deep Belief Nets (DBNs)
- Denoising Autoencoders
- Stacked Denoising Autoencoders (SDA)
- **Summary**

# Summary

- The problem of the previous neural networks algorithm
- The breakthrough of deep learning
- RBM, Restricted Boltzmann Machines
- DBN with RBM
- DA, Denoising Autoencoders
- SDA with DA